EVALFORGE – LLM Evaluation & Hallucination Detection Platform

OBJECTIVE:

Build a production-grade system to evaluate and benchmark LLM outputs, detect hallucinations, and improve reliability using RAG and semantic scoring.

PROBLEM:

LLMs produce hallucinations and unreliable outputs. No standardized evaluation system exists.

CORE FEATURES:

- LLM Evaluation Engine (semantic similarity, relevance, hallucination score)

- Hallucination Detection using embeddings + grounding

- Multi-model benchmarking (Gemini, OpenAI)

- RAG evaluation (retrieval quality, grounding score)

- Web dashboard for visualization

DIFFERENTIATORS:

1. Demo Video (system walkthrough)

2. Architecture Diagram (system design clarity)

3. Real Dataset (SQuAD / TruthfulQA)

4. Metrics Dashboard (accuracy, latency, hallucination rate)

TECH STACK:

Backend: FastAPI, LangChain

ML: Sentence Transformers, HuggingFace

Frontend: React, Tailwind

Infra: Docker, Redis, Vercel

SUCCESS METRICS:

- 85% accuracy

- 70% hallucination detection precision

- <300ms latency

- 1000+ dataset samples

GITHUB STRUCTURE:

backend/, frontend/, datasets/, notebooks/, architecture.png, demo.gif

RESUME IMPACT:

- Built LLM evaluation system
- Reduced hallucination
- Benchmarked models
- Designed scalable architecture